

STRATEGIC BRIEFING | INTERNAL

Agent-Ready 1325.AI

The MCP Advantage: A Rare Position in an Emerging Category

Prepared for: The 1325.AI Executive Team

Prepared by: Thomas D. Bowling, Founder & Chief Architect

Date: July 08, 2026

Company founded: June 5, 2025

Document reference: 1325AI-BRF-2026-003 (v3)

A Note from the Founder

In thirteen months — from June 5, 2025 to today — 1325.AI has moved from a directory concept to a live, agent-addressable platform sitting on the same protocol rail that ChatGPT, Claude, Cursor, and every serious enterprise assistant now use to reach the outside world. That is not a small accomplishment; it is a strategic position most companies in our category will not hold for another thirteen to thirty-six months, if they get there at all.

This briefing is written for the team that built it and the partners who will carry it forward. It explains, in plain terms, what the Model Context Protocol is, why our early implementation is structurally rare, and how we intend to convert that rarity into durable commercial advantage.

— Thomas D. Bowling, Founder & Chief Architect

Executive Summary

1325.AI is now discoverable and callable by autonomous AI agents through the Model Context Protocol (MCP) — the open standard, introduced by Anthropic in late 2024 and adopted across the industry in 2025–2026, that lets AI assistants securely call external tools and data. Our directory of 43,000+ verified businesses, loyalty rewards, and per-user QR scan history is exposed as a set of five read-only, OAuth-secured tools that any compliant AI client can consume with no app install, no scraping, and no bespoke integration work.

The strategic implication is a category shift: 1325.AI is no longer competing only as a consumer marketplace. It is now positioned as **agent-grade infrastructure** — the type of asset that is priced and acquired on infrastructure multiples, not marketplace multiples. This briefing details the mechanics, the rarity, the risks we are actively managing, and the actions required of the team.

1. What the Model Context Protocol Actually Is

MCP is an open, JSON-RPC-based protocol that standardizes how a Large Language Model (LLM) assistant discovers, authenticates against, and invokes external tools. Introduced by Anthropic in November 2024 and rapidly adopted by OpenAI, Google, Microsoft, and the major IDE and agent vendors through 2025, it now functions as the de facto integration layer for agentic AI — the same role that HTTP played for the early web and that REST plus OAuth 2.0 played for the SaaS era.

Three properties matter for our position:

- 1. Discovery is machine-native.** An MCP server publishes a manifest listing its tools, input schemas, and auth requirements. Agents introspect it automatically — no human integration engineer required. When a user asks ChatGPT or Claude to 'find a Black-owned electrician in Atlanta,' our *search_directory* tool is a candidate the model can select on its own.
- 2. Auth is enterprise-grade by default.** MCP prescribes OAuth 2.1 with Dynamic Client Registration. Our implementation runs on Supabase as the authorization server; the AI client connects as *the end user*, and every database call is filtered by Row-Level Security under that user's identity. There is no shared API key, no service-role escalation, no bespoke session handling. This is the security posture Fortune-500 IT organizations require before they will allow an agent to touch a third-party system.
- 3. Distribution is zero-friction.** The user does not install our app, sign up separately, or copy an API key. They add 1325.AI as a connector inside the assistant they already use, approve consent once, and every future question they ask that touches our domain can be answered with our live data. This collapses the customer-acquisition funnel from five steps to one.

2. Why This Position Is Structurally Rare

The 2026 MCP Landscape — In Numbers	
Public MCP servers indexed globally (est. Q2 2026)	~1,800
Of those, servers with OAuth 2.1 + DCR (not static tokens)	< 12%
Servers backed by a verified domain-specific dataset > 25,000 records	< 3%
Servers combining all three <i>and</i> a defensible category niche	a low double-digit count
1325.AI satisfies all four conditions today.	

The rarity is not that we built an MCP server — anyone can. The rarity is the **combination**: (a) a production-grade OAuth implementation that Fortune-500 procurement can approve without modification; (b) a proprietary, verified dataset of scale that cannot be reconstructed by scraping; (c) a live user base generating first-party signal (scans, redemptions, points) that improves the underlying model over time; and (d) a category — verified Black-owned commerce — with no incumbent MCP presence and clear cultural and civic demand.

In practical terms, the closest historical analogue is a small set of companies that shipped public REST APIs in 2006–2008 (Twilio, Stripe, early Twitter) before the category was named. Their early protocol position translated into pricing power, distribution partnerships, and eventual acquisition or IPO on multiples that had no relationship to their 2006 revenue. We are entering the same window on the agent rail.

3. How This Benefits 1325.AI, In Detail

A. Category re-rating from marketplace to infrastructure.

Consumer marketplaces are valued on gross merchandise volume (GMV) and take-rate — typically 3×–6× revenue at maturity. API and protocol infrastructure businesses are valued on usage growth, developer adoption, and lock-in — commonly 15×–40× revenue in the 2023–2026 comparable set (Twilio, Stripe, Cloudflare, MongoDB at IPO stage). Even a partial re-rating materially changes our enterprise value and our leverage in every capital conversation.

B. Distribution surface area, without a distribution cost.

Every user of ChatGPT, Claude, Cursor, Perplexity, Microsoft Copilot, and the next wave of vertical agents is a latent user of 1325.AI. We did not have to build that reach, we did not have to pay for the app-store slot, and we do not carry the customer-acquisition cost line-item that traditional directories carry. The economics of a marginal agent-driven query are near zero for us and highly valuable to the agent vendor — which is exactly the asymmetry that produces revenue-share and preferred-partner deals.

C. A defensible data flywheel.

Every agent-mediated query produces a structured, consented event: what was asked, what was returned, what the user did next. That signal — not the directory itself — is the long-term moat. Competitors who ship an MCP server in 2027 will not have the 18-month behavioral corpus we will have accumulated by then.

D. Enterprise and civic sales acceleration.

For denominational, municipal, and Fortune-500 partners, 'we are already an MCP-native provider with OAuth 2.1' is a procurement-ready sentence. It compresses the security-review cycle from months to weeks and moves

us from 'directory vendor' to 'infrastructure vendor' in the conversation — which is a different budget line and a different decision-maker.

E. Patent and IP posture strengthened.

U.S. Provisional Patent Application No. 63/969,202 (27 claims pending) covers the scoring, temporal founding-status, and per-user QR-loyalty mechanics that sit *behind* the MCP surface. Exposing read-only tools while keeping the write paths, scoring formulas, and settlement logic proprietary is a deliberate choice: it maximizes distribution while preserving the patentable core.

F. Optionality on the next protocol wave.

MCP is the current standard. Whatever succeeds or extends it — agent-to-agent negotiation, payment-capable tool calls, verifiable-credential handoffs — will inherit MCP's shape. Teams that already operate a production MCP server are the natural first adopters of every successor specification. We are buying an option on the next five years of protocol evolution.

4. What We Deliberately Held Back (and Why)

The v1 MCP surface exposes five read-only tools: *search_directory*, *get_business*, *list_rewards*, *get_my_points_balance*, and *get_my_recent_scans*. There is no payment tool, no QR write path, no scoring endpoint, and no reward-redemption tool. This is not a limitation; it is a strategy with three payoffs:

- Compliance surface stays narrow.** Read-only tools do not trigger Apple StoreKit, Stripe connected-account, or PCI review. We ship faster and defend more easily.
- Patentable mechanics remain private.** The scoring formula, the atomic *award_qr_scan* RPC, and the temporal founding-status logic are the differentiated IP — they stay server-side, under RLS, and under patent.
- Read-only is the correct stress test.** We learn how agents actually query our data — vocabulary, latency, error rates — before we open transactional surface area. Every serious infrastructure company (Stripe, Twilio, Plaid) shipped read endpoints before write endpoints.

5. Execution Risks We Are Actively Managing

Risk	Mitigation
Edge Function cold-start latency degrades agent UX	Warm-path pinging, response caching for high-frequency queries, and a p95 latency SLO tracked weekly.
Semantic tuning of <i>search_directory</i> (agents phrase queries differently than humans)	Query-log review cadence with prompt-level adjustments; embedding-based ranking is on the roadmap once query volume justifies it.
OAuth consent-screen abandonment	Consent copy is A/B tested; drop-off is instrumented at each step of the authorization flow.
Fast-follower risk from larger directories	Data verification, patent claims, and first-party behavioral signal are the durable moat; protocol presence alone is not.
Agent-vendor policy shifts (e.g. connector-store gating)	Multi-client posture from day one: ChatGPT, Claude, Cursor, and any MCP-compliant client. No single-vendor dependency.

6. What This Means for the Team

Every function of the company should update its posture:

- Sales & Partnerships:** lead with 'agent-native infrastructure,' not 'directory.' Every deck, one-pager, and outbound email should reference MCP, OAuth 2.1, and the 43,000-business corpus in the first two sentences.
- Engineering:** treat the MCP surface as a product with its own SLOs, changelog, and deprecation policy. Backwards compatibility on tool schemas is now a customer promise.
- Data & Operations:** verification quality is now visible to every AI assistant on the market. A bad address in our database is a bad answer in ChatGPT. Verification cadence moves from 'quarterly' to 'continuous.'
- Investor & Capital:** the valuation narrative shifts from $GMV \times \text{take-rate}$ to agent-mediated query volume, connected-client count, and revenue per active connector.
- Legal & IP:** continue prosecution of 63/969,202 on schedule; begin scoping the non-provisional and any continuations that read on MCP-mediated commerce mechanics.

7. Five-Step Activation Plan

1. Publish	Confirm the MCP endpoint is live on production and referenced from /connect. Verify OAuth flow end-to-end from a clean ChatGPT and Claude account.
2. Register	Submit the connector to the ChatGPT and Claude developer consoles. Track listing status weekly.
3. Instrument	Wire per-tool telemetry: invocation count, p50/p95 latency, error rate, and downstream conversion (scan, sign-up, redemption).
4. Reposition	Refresh the top of every sales, sponsor, and investor deck to lead with agent-native infrastructure. Retire language that frames us as a directory only.
5. Outreach	Coordinated announcement to denominational, municipal, and Fortune-500 partners with a procurement-ready one-pager attached.

Closing Note

We built this in thirteen months. The window to be one of the first credible, OAuth-secured, dataset-backed MCP servers in our category will not stay open indefinitely — the same protocol adoption curve that created our opening is what will eventually close it. The task in front of the team is not to celebrate the position; it is to convert it, quickly and deliberately, into partnerships, revenue, and defensible market share while the advantage is still asymmetric.

— Thomas D. Bowling, Founder & Chief Architect